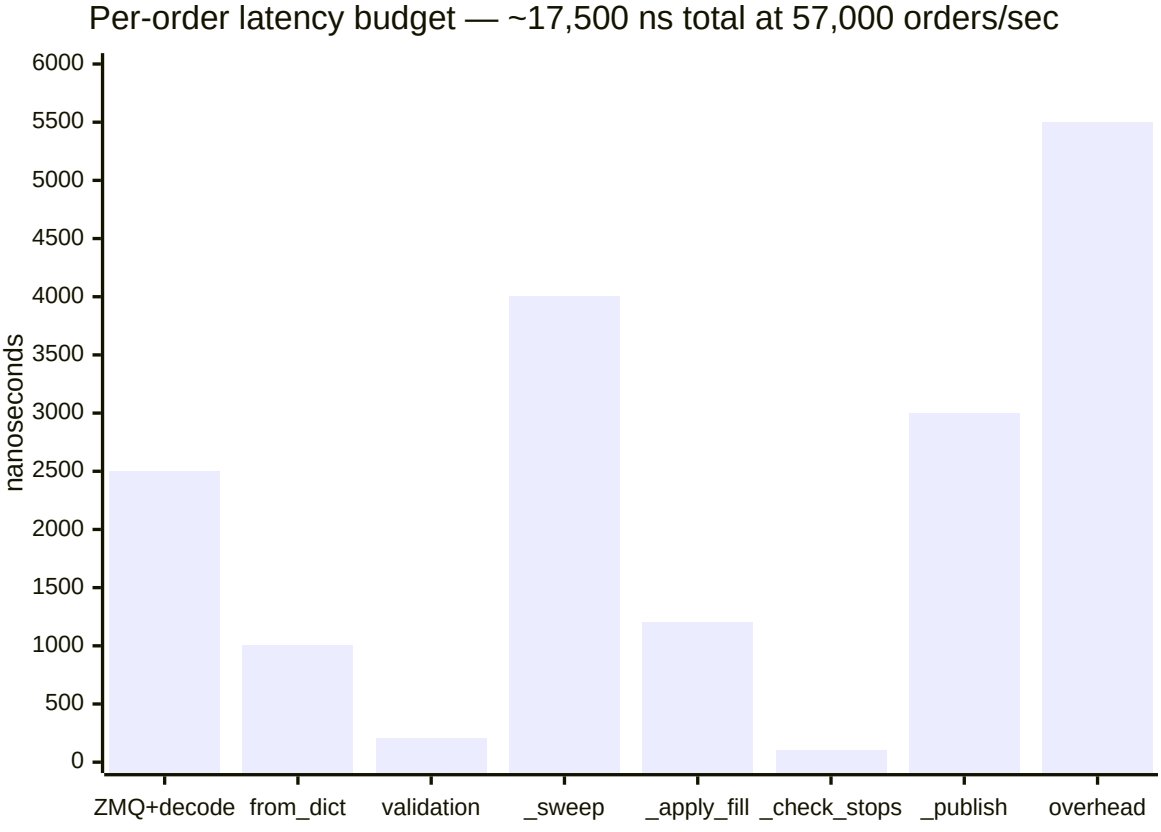
Per-order latency budget — ~17,500 ns total at 57,000 orders/sec
nanoseconds
6000
5500
5000
4500
4000
3500
3000
2500
2000
1500
1000
500
0
ZMQ+decode
from_dict
validation
_sweep
_apply_fill
_check_stops
_publish
overhead